


How a question becomes a cited answer

The whole system on one page: what happens once per repository, what happens on every question, and the three stores that connect the two. Line colour encodes where the work runs.



Two lanes, one shared middle. Lane A runs once and fills the stores; lane C runs on every question and reads them. Nothing in lane C touches the repository — by then the code has already been parsed, embedded and indexed, which is what makes answering fast.

— runs on your machine — network call — OpenAI, Cohere, Qdrant Cloud ■■■ local disk — lost on container restart

Reading it in order

1. **Ingestion (A)**. You paste a GitHub URL. The API saves a row as `pending` and returns immediately — the real work takes minutes, so it runs in the background. The repo is shallow-cloned, vendored and generated files are dropped, and tree-sitter cuts what's left at function and class boundaries. Each chunk gets embedded twice — a dense vector from OpenAI for meaning, a sparse BM25 vector computed locally for exact keywords — then upserted. The call graph is saved alongside, and the row flips to `ready`.
2. **Stores (B)**. Three of them, each holding what it is good at: vectors in Qdrant, relational data in SQLite, and the call graph as JSON because it is read whole and written whole.
3. **Query (C)**. Your question is rewritten into something that matches code rather than conversation. Both indexes are searched and the two rankings are fused by position, not score. A cross-encoder reranks the survivors and — crucially — reports how confident it is. If the question is about how code *flows*, the call graph pulls in callers and callees; otherwise that step is skipped, because extra context buries the answer to a narrow question.
4. **Outcome (D)**. The answer must cite `path:line`. If confidence was low, or the model produced prose with no citation, the system refuses instead of guessing. Either way the result is saved and streamed back. The next question reuses the same index — ingestion never runs again unless you re-add the repo.